

# Module 6

## Active database

## Deductive database

## Temporal database

### Outline

- Active database & triggers
- Temporal databases
- Deductive databases

Dr. Geetha Mary A  
Associate Professor,  
SCOPE, Vellore Institute of Technology,  
Vellore

Source:  
Pearson Education, Inc. 2011. Elmasri/Navathe, Fundamentals of Database Systems, Sixth Edition

Slide - 2

## Active Database Concepts and Triggers

### Generalized Model for Active Databases and Oracle Triggers

- **Triggers** are executed when a specified condition occurs during insert/delete/update
- Triggers are action that fire automatically based on these conditions

## Event-Condition-Action (ECA) Model

### Generalized Model (contd.)

- Triggers follow an Event-condition-action (ECA) model
  - **Event:**
    - Database modification
    - E.g., insert, delete, update,
  - **Condition:**
    - Any true/false expression
    - Optional: If no condition is specified then condition is always true
  - **Action:**
    - Sequence of SQL statements that will be automatically executed

Slide - 3

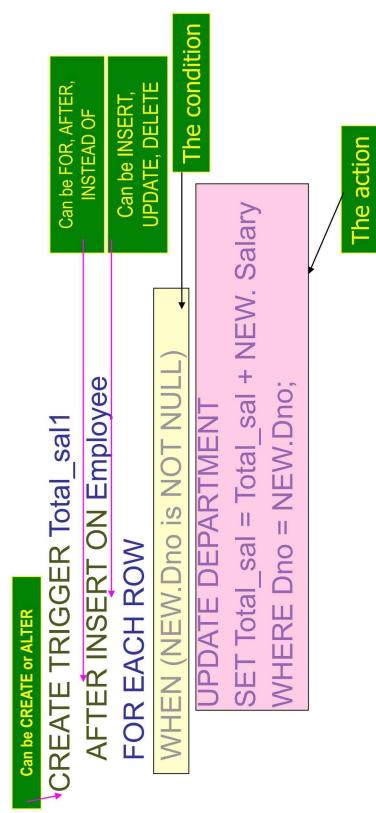
Slide - 4

## Trigger Example

### Generalized Model (contd.)

- When a new employee is added to a department, modify the Total\_sal of the Department to include the new employees salary
- Logically this means that we will CREATE a TRIGGER, let us call the trigger Total\_sal1
  - Condition
    - This trigger will execute AFTER INSERT ON Employee table
    - It will do the following FOR EACH ROW
      - WHEN NEW.Dno is NOT NULL
      - The trigger will UPDATE DEPARTMENT
      - By SETting the new Total\_sal to be the sum of
        - old Total\_sal and NEW. Salary
        - WHERE the Dno matches the NEW.Dno;

Slide - 5



Slide - 6

## Example: Trigger Definition

## CREATE or ALTER TRIGGER

### Generalized Model (contd.)

- CREATE TRIGGER <name>
  - Creates a trigger
- ALTER TRIGGER <name>
  - Alters a trigger (assuming one exists)
- CREATE OR ALTER TRIGGER <name>
  - Creates a trigger if one does not exist
  - Alters a trigger if one does exist
  - Works in both cases, whether a trigger exists or not

### Generalized Model (contd.)

- AFTER
  - Executes after the event
- BEFORE
  - Executes before the event
- INSTEAD OF
  - Executes instead of the event
    - Note that event does not execute in this case
      - E.g., used for modifying views

## Conditions

### Generalized Model (contd.)

- CREATE TRIGGER <name>
  - Creates a trigger
- ALTER TRIGGER <name>
  - Alters a trigger (assuming one exists)
- CREATE OR ALTER TRIGGER <name>
  - Creates a trigger if one does not exist
  - Alters a trigger if one does exist
  - Works in both cases, whether a trigger exists or not

Slide - 7

Slide - 8

## Row-Level versus Statement-level

### Condition

#### Generalized Model (contd.)

- Triggers can be
  - **Row-level**
    - FOR EACH ROW specifies a row-level trigger
  - **Statement-level**
    - Default (when FOR EACH ROW is not specified)
    - Row level triggers
      - Executed separately for each affected row
    - Statement-level triggers
      - Execute once for the SQL statement,

Slide - 9

Slide - 10

### Action

#### Generalized Model (contd.)

- Action can be
  - One SQL statement
  - A sequence of SQL statements enclosed between a BEGIN and an END

- Action specifies the relevant modifications

### Triggers on Views

#### Generalized Model (contd.)

- INSTEAD OF triggers are used to process view modifications

Slide - 11

Slide - 12

## Active Database Concepts and Triggers

- Design and Implementation Issues for Active Databases
  - An active database allows users to make the following changes to triggers (rules)
    - Activate
    - Deactivate
    - Drop

Slide - 13

## Active Database Concepts and Triggers

- Design and Implementation Issues for Active Databases
  - An event can be considered in 3 ways
    - Immediate consideration
    - Deferred consideration
    - Detached consideration

Slide - 14

## Active Database Concepts and Triggers

- Design and Implementation Issues (contd.)
  - Immediate consideration
  - Part of the same transaction and can be one of the following depending on the situation
    - Before
    - After
    - Instead of
  - Deferred consideration
    - Condition is evaluated at the end of the transaction
    - Detached consideration
    - Condition is evaluated in a separate transaction
- Potential Applications for Active Databases
  - Notification
    - Automatic notification when certain condition occurs
    - Enforcing integrity constraints
    - Triggers are smarter and more powerful than constraints
  - Maintenance of derived data
    - Automatically update derived data and avoid anomalies due to redundancy
      - E.g., trigger to update the Total\_Sal in the earlier example

Slide - 15

Slide - 16

## Active Database Concepts and Triggers

### Triggers in SQL-99

- Can alias variables inside the REFERENCING clause

### ■ Trigger examples

```
T1: CREATE TRIGGER Total_sal1  
AFTER UPDATE OF Salary ON EMPLOYEE  
REFERENCING OLD ROW AS O, NEW ROW AS N  
FOR EACH ROW  
WHEN ( NDno IS NOT NULL )  
UPDATE DEPARTMENT  
SET Total_sal = Total_sal + N.salary - O.salary  
WHERE Dno = N.Dno;  
  
T2: CREATE TRIGGER Total_sal2  
AFTER UPDATE OF Salary ON EMPLOYEE  
REFERENCING OLD TABLE AS O, NEW TABLE AS N  
FOR EACH STATEMENT  
WHEN EXISTS ( SELECT * FROM N WHERE N.Dno IS NOT NULL ) OR  
EXISTS ( SELECT * FROM O WHERE O.Dno IS NOT NULL )  
UPDATE DEPARTMENT AS D  
SET D.Total_sal = D.Total_sal  
+ ( SELECT SUM (N.Salary) FROM N WHERE D.Dno=N.Dno )  
- ( SELECT SUM (O.Salary) FROM O WHERE D.Dno=O.Dno )  
WHERE Dno IN ( SELECT Dno FROM N ) UNION ( SELECT Dno FROM O );
```

Slide - 17

Slide - 18

## Active Database Concepts and Triggers

### Temporal Database Concepts

- Temporal databases require some aspect of time when organizing information

- Healthcare
- Insurance
- Reservation systems
- Scientific databases

SQL2 temporal data types

- DATE, TIME, TIMESTAMP, INTERVAL, PERIOD

## Temporal Database Concepts

- Time Representation, Calendars, and Time Dimensions
- Time is considered ordered sequence of points in some granularity
- Use the term chronon instead of point to describe minimum granularity

2) A temporal database stores data relating to time instances. It offers temporal data types and stores information relating to past, present and future time. Temporal databases could be uni-temporal or tri-temporal.

Uni-Temporal [ edit ]

A uni-temporal database has one axis of time, either the validity range or the system time range.

Bi-Temporal [ edit ]

A bi-temporal database has two axes of time:

- valid time
- transaction time or decision time

A tri-temporal database has three axes of time:

- valid time
- transaction time
- decision time

4)

Slide - 19

## Temporal Database Concepts

### Time Representation, ... (contd.)

- A **calendar** organizes time into different time units for convenience.

- Accommodates various calendars

Gregorian (western)

- Chinese ✓ Valid time is the time period during which a fact is true in the real world.
- Islamic ✓ Transaction time is the time at which a fact was recorded in the database.
- Hindu ✓ Decision time is the time at which the decision was made about the fact.
- Jewish ✓
- Etc.

Slide - 20

## Temporal Database Concepts

### Time Representation, ... (contd.)

#### Point events

- Single time point event
  - E.g., bank deposit

- Series of point events can form a **time series data**

#### Duration events

- Associated with specific time period
  - Time period is represented by **start time** and **end time**

## Temporal Database Concepts

### Time Representation, ... (contd.)

#### Transaction time

- The time when the information from a certain transaction becomes valid

#### Bitemporal database

- Databases dealing with two time dimensions

## Features [ edit ]

Temporal databases support managing and accessing temporal data by providing one or more of the following features:[1][2]

- A time period datatype, including the ability to represent time periods with no end (infinity or forever)
  - The ability to define valid and transaction time period attributes and bitemporal relations
  - System-maintained transaction time
- Temporal primary keys, including non-overlapping period constraints
- Temporal constraints, including non-overlapping uniqueness and referential integrity
- Update and deletion of temporal records with automatic splitting and coalescing of time periods
- Temporal queries at current time, time points in the past or future, or over durations
- Predicates for querying time periods, often based on Allen's interval relations

Slide - 21

Slide - 22

## Temporal Database Concepts

### Incorporating Time in Relational Databases Using Tuple Versioning

- Add to every tuple
  - Valid start time
  - Valid end time

## Temporal Database Concepts

**Figure 24.7**  
Different types of temporal relational databases.  
(a) Valid time database schema.  
(b) Transaction time database schema.  
(c) Bitemporal database schema.

VALID TIME data-base schema		(a) EMP_VT			
DEPT_VT		Name	Ssn	Salary	Dno
		Dname	Dno	Total_sal	Manager_ssn
				Vst	Vet

TRANSACTION TIME data-base schema		(b) EMP_TT			
DEPT_TT		Name	Ssn	Salary	Dno
		Dname	Dno	Total_sal	Manager_ssn
				Tst	Tet

BITEMPORAL data-base schema		(c) EMP_BT			
DEPT_BT		Name	Ssn	Salary	Dno
		Dname	Dno	Total_sal	Manager_ssn
				Tst	Tet

Slide - 23

Slide - 24

## Temporal Database Concepts

### Incorporating Time in Object-Oriented Databases Using Attribute Versioning

- A single complex object stores all temporal changes of the object
- **Time varying attribute**
  - An attribute that changes over time
    - E.g., age ✓
  - **Non-Time varying attribute**
    - An attribute that does **not** change over time
      - E.g., date of birth ✓

**Figure 24.8**  
Some tuple versions in the valid time relations EMP\_VT and DEPT\_VT.

EMP_VT					
Name	Ssn	Salary	Dno	Supervisor_ssn	Vst
Smith	123456789	25000	5	333445555	2002-06-15
Smith	123456789	30000	5	333445555	2003-06-01
Wong	333445555	25000	4	999887777	1999-08-20
Wong	333445555	30000	5	999887777	2001-02-01
Wong	333445555	40000	5	888665555	2002-04-01
Brown	222447777	28000	4	999887777	2001-05-01
Narayan	666884444	38000	5	333445555	2003-08-01
...					
DEPT_VT					
Dname	Dno	Manager_ssn	Vst	Vet	
Research	5	888665555	2001-09-20	2002-03-31	
Research	5	333445555	2002-04-01	Now	

Slide - 25

Slide - 26

- **Types of updates**

- **Proactive**
- **Retroactive**
- **Simultaneous**

- Implementation considerations
  - Store all tuples in the same table
  - Create two tables: one for currently valid information and one for the rest
  - Vertically partition temporal relation attributes into separate relations
    - New tuple created whenever any attribute updated
  - Append-only database
  - Keeps complete record of changes and corrections

```

class TEMPORAL_SALARY
{
  attribute Date           Valid_start_time;
  attribute Date           Valid_end_time;
  attribute float          Salary;
};

class TEMPORAL_DEPT
{
  attribute Date           Valid_start_time;
  attribute Date           Valid_end_time;
  attribute DEPARTMENT_VT Dept;
};

class TEMPORAL_SUPERVISOR
{
  attribute Date           Valid_start_time;
  attribute Date           Valid_end_time;
  attribute EMPLOYEE_VT   Supervisor;
};

class TEMPORAL_LIFESPAN
{
  attribute Date           Valid_start_time;
  attribute Date           Valid_end_time;
};

class EMPLOYEE_VT
{
  extent EMPLOYEES
  {
    attribute list<TEMPORAL_LIFESPAN> Lifespan;
    attribute string Name;
    attribute string Ssn;
    attribute string SalHistory;
    attribute string DeptHistory;
    attribute list<TEMPORAL_DEPT> SupervisorHistory;
  };
};

```

- **CREATE TABLE statement**

- Extended with optional AS clause
- Allows users to declare different temporal options

- **Examples:**

- AS VALID STATE<GRANULARITY> (valid time relation with valid time period)
- AS TRANSACTION (transaction time relation with transaction time period)
- **Keywords STATE and EVENT**
  - Specify whether a time period or point associated with valid time dimension

- Time series data
  - Often used in financial, sales, and economics applications
  - Special type of valid event data
  - Event's time points predetermined according to fixed calendar
    - Managed using specialized time series management systems
    - Supported by some commercial DBMS packages

## Introduction to Deductive Databases

- Overview of Deductive Databases
- Prolog/Datalog Notation
- Datalog Notation
- Clausal Form and Horn Clauses
  - Interpretation of Rules
  - Datalog Programs and Their Safety
  - Use the Relational Operations
  - Evaluation of Non-recursive Datalog Queries

### What is a deductive database system?

A deductive database can be defined as an advanced database augmented with an inference system.



By evaluating rules against facts, new facts can be derived, which in turn can be used to answer queries. It makes a database system more powerful.

## Overview of Deductive Databases

- **Declarative Language**
- **Language to specify rules**
- **Inference Engine (Deduction Machine)**
- Can deduce new facts by interpreting the rules
- Related to logic programming
  - Prolog language (Prolog => Programming in logic)
- Uses backward chaining to evaluate
  - Top-down application of the rules

In a deductive database system we typically specify rules through a **declarative language**—a language in which we specify what to achieve rather than how to achieve it. An **inference engine** (or **deduction mechanism**) within the system can deduce new facts from the database by interpreting these rules. The model used for deductive databases is closely related to the relational data model and particularly to the domain relational calculus formalism (see Section 6.6). It is also related to the field of **logic programming** and the **Prolog** language. The deductive database work based on logic has used Prolog as a starting point. A variation of Prolog called **Datalog** is used to define rules declaratively in conjunction with an existing set of relations, which are themselves treated as literals in the language. Although the language structure of Datalog resembles that of Prolog, its operational semantics—that is, how a Datalog program is executed—is still different.

**Facts are specified in a manner similar to the way relations are specified**, except that it is not necessary to include the attribute names. Recall that a tuple in a relation describes some real-world fact whose meaning is partly determined by the attribute names. In a deductive database, the meaning of an attribute value in a tuple is determined solely by its **position** within the tuple.

**Rules are somewhat similar to relational views**. They specify virtual relations that are not actually stored but that can be formed from the facts by applying inference mechanisms based on the rule specifications. The main difference between rules and views is that rules may involve **recursion** and hence may yield virtual relations that cannot be defined in terms of basic relational views.

*Rules → View*

## Overview of Deductive Databases

- Speciation consists of:
  - Facts
  - Similar to relation specification without the necessity of including attribute names
- Rules

### Prolog/Datalog Notation

Example (a logic program)

```
Facts:
supervise(franklin, john),
supervise(franklin, ramesh),
supervise(franklin, joyce),
supervise(james, franklin),
supervise(jennifer, alicia),
supervise(jennifer, jeffrey),
supervise(james, jennifer).

Rules:
superior(X, Y) :- supervises(X, Y),
superior(X, Y) :- supervises(X, Z), superior(Z, Y),
superior(X, Y) :- supervises(X, Y),
superior(Y, X) :- superior(X, Y).

subordinate(X, Y) :- superior(Y, X).
```

Slide - 29

### Prolog/Datalog Notation

Diagram illustrating the structure of the logic program:

```
graph TD
    James --- Franklin
    James --- Jennifer
    Franklin --- John
    Franklin --- Ramesh
    Franklin --- Joyce
    Jennifer --- Alicia
    Jennifer --- Jeffery
    James --- Jennifer

    Superior[superior(X, Y)] --> SupervisesX[supervises(X, Y)]
    Superior --> SuperiorXY[superior(X, Y) :- supervises(X, Z), superior(Z, Y)]
    Superior --> SuperiorYX[superior(Y, X) :- superior(X, Y)]

    Subordinate[X, Y] --> SuperiorYX[superior(Y, X) :- superior(X, Y)]
```

Slide - 30

## DATALOG

### Prolog/Datalog Notation

- SQL queries can be read as follows:
  - If some tuples exist in the From tables that satisfy the Where conditions, then the Select tuple is in the answer."
- Datalog is a query language that has the same if,then flavor:
  - New: The answer table can appear in the From clause, i.e., be defined recursively.
- Predicate has
  - a name
  - a fixed number of arguments
- Convention:
  - Prolog style syntax is commonly used.
  - Constants are numeric or character strings
  - Variables start with upper case letters
- E.g., SUPERVISE(Supervisor, Supervisee)
  - States that Supervisor SUPERVISE(s) Supervisee

- Example (a logic program)
- Facts:
  - supervise(franklin, john),  
supervise(franklin, ramesh),  
supervise(franklin, joyce),  
supervise(james, franklin),  
supervise(jennifer, alicia),  
supervise(jennifer, jeffrey),  
supervise(james, jennifer).
- Rules:
  - superior(X, Y) :- supervises(X, Y),  
superior(X, Y) :- supervises(X, Z), superior(Z, Y),  
superior(X, Y) :- supervises(X, Y),  
superior(Y, X) :- superior(X, Y).
  - subordinate(X, Y) :- superior(Y, X).

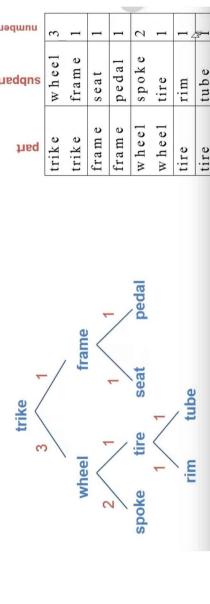
- E.g., SUPERVISE(Supervisor, Supervisee)
  - States that Supervisor SUPERVISE(s) Supervisee

### Prolog/Datalog Notation

#### Rule

- Is of the form head :- body
  - where :- is read as if and only iff
- E.g., SUPERIOR(X,Y) :- SUPERVISE(X,Y)
- E.g., SUBORDINATE(Y,X) :- SUPERVISE(X,Y)

Example

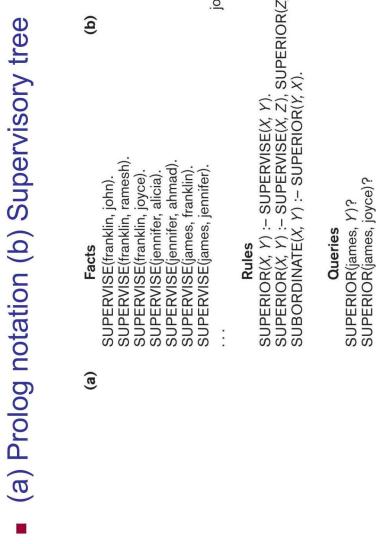


#### Query

- Involves a predicate symbol followed by some variable arguments
- A Datalog Query to answer the question
  - where :- is read as if and only iff
- E.g., SUPERIOR(james, Y)?
  - Comp(Part, Subpt) :- Assembly(Part, Subpt, Qty).  
Comp(Part, Subpt) :- Assembly(Part, Part2, Qty).  
Comp(Part2, Subpt).  
head of rule      implication  
body of rule
- E.g., SUBORDINATE(james,X)?
  - Can read the second rule as follows:  
"For all values of Part, Subpt and Qty,  
if there is a tuple (Part, Part2, Qty) in Assembly  
and a tuple (Part2, Subpt) in Comp,  
then there must be a tuple (Part, Subpt) in Comp."

Slide - 31

**Figure 24.11**



## Datalog Notation

- Datalog notation
  - Program is built from **atomic formulae**
  - **Literals** of the form  $p(a_1, a_2, \dots, a_n)$  where
    - $p$  predicate name
    - $n$  is the number of arguments
    - Built-in predicates are included
    - E.g.,  $<$ ,  $\leq$ , etc.
  - **A literal** is either
    - An atomic formula
    - An atomic formula preceded by not

Slide - 33

Slide - 34

## Clausal Form and Horn Clauses

- A formula can have quantifiers
  - Universal
  - Existential

## Clausal Form and Horn Clauses

- In clausal form, a formula must be transformed into another formula with the following characteristics
  - All variables are universally quantified
  - Formula is made of a number of clauses where each clause is made up of literals connected by logical ORs only
  - Clauses themselves are connected by logical ANDs only.

Slide - 35

Slide - 36

## Clausal Form and Horn Clauses

- Any formula can be converted into a clausal form
  - A specialized case of clausal form are horn clauses that can contain no more than one positive literal
- Datalog program are made up of horn clauses

- Clause

A	B	$\neg A \vee B$
1	1	1
0	1	1
1	0	0
0	0	1

Horn clause  
A Horn clause is a clause with the head containing only one positive atom.  
$$\underline{\quad \quad \quad B_m \leftarrow A_1, \dots, A_n}$$

Slide - 37

Slide - 38

## Interpretation of Rules

- There are two main alternatives for interpreting rules:
  - Proof-theoretic
  - Model-theoretic
- Datalog program are made up of horn clauses

A	B	$\neg A \vee B$
1	1	1
0	1	1
1	0	0
0	0	1

Horn clause  
A Horn clause is a clause with the head containing only one positive atom.  
$$\underline{\quad \quad \quad B_m \leftarrow A_1, \dots, A_n}$$

Slide - 38

Slide - 39

## Interpretation of Rules

- Proof-theoretic
- Facts and rules are **axioms**
- **Ground axioms** contain no variables
- Rules are **deductive axioms**
- **Deductive axioms** can be used to construct new facts from existing facts
  - This process is known as theorem proving

1. SUPERIOR(X, Y) :- SUPERVISE(X, Y).  
(rule 1)
2. SUPERIOR(X, Y) :- SUPERVISE(X, Z), SUPERIOR(Z, Y).  
(rule 2)
3. SUPERVISE(jennifer, ahmad).
4. SUPERVISE(james, jennifer).
5. SUPERIOR(jennifer, ahmad).
6. SUPERIOR(james, ahmad).

Slide - 39

Slide - 40



## Datalog Programs and Their Safety

- Two main methods of defining truth values
  - Fact-defined predicates (or relations)
    - Listing all combination of values that make a predicate true
  - Rule-defined predicates (or views)
    - Head (LHS) of 1 or more Datalog rules, for example Figure 24.15

```
SUPERIOR(X, Y) :- SUPERVISOR(X, Y).
SUPERIOR(X, Y) :- SUPERVISOR(X, Z), SUPERIOR(Z, Y).

SUBORDINATE(X, Y) :- SUPERIOR(Y, X).

SUPERVISOR(X) :- EMPLOYEE(X), SUPERVISOR(X, Y).

OVER_40K_EMP(X) :- EMPLOYEE(X), SALARY(X, Y), Y >= 40000.
UNDER_40K_SUPERVISOR(X) :- SUPERVISOR(X), NOT(OVER_40K_EMP(X)).
MAIN_PRODUCT(X, EMPLOYEE(X)) :- EMPLOYEE(X), WORKS_ON(X, productX, Y), Y >= 20.
PRESIDENT(X) :- EMPLOYEE(X), NOT(SUPERVISOR(X, Y)).
```

Slide - 45

## Datalog Programs and Their Safety

- A program is **safe** if it generates a **finite set of facts**
  - Fact-defined predicates (or relations)
  - Rule-defined predicates (or views)
- Head (LHS) of 1 or more Datalog rules, for example Figure 24.15

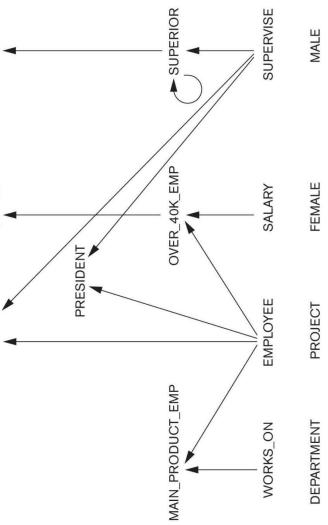
Slide - 46

## Use the Relational Operations

- Many operations of relational algebra can be defined in the form of Datalog rules that define the result of applying these operations on database relations (fact predicates)
  - Relational queries and views can be easily specified in Datalog

## Evaluation of Non-recursive Datalog Queries

- Define an inference mechanism based on relational database query processing concepts
  - See Figure 24.17 on predicate dependencies for Figs 24.14 and 24.15



Slide - 47

Slide - 48