

Socket Programming

Data types and structures for writing client-server programs

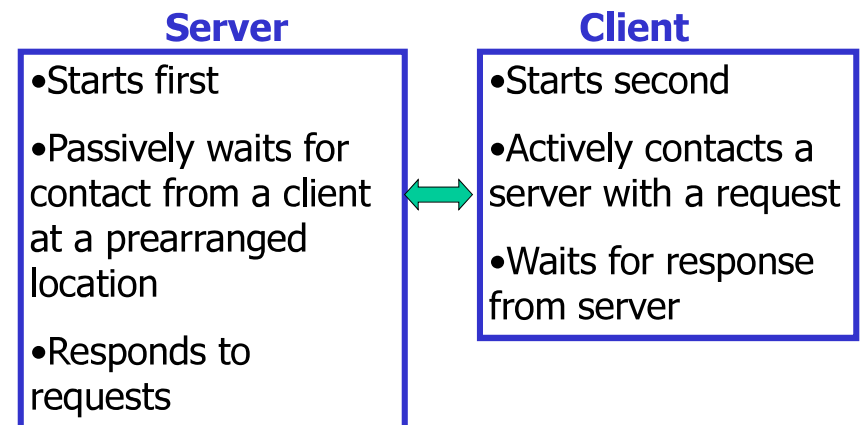
Introduction

- Transport layer and layers below
 - Basic communication
 - reliability
- Application Layer Functionality
 - Abstractions
 - Names:
 - define symbolic names to identify both physical and abstract resources available on an internet

- Network
 - transfers bits
 - operates at application's request
- Application determines
 - what/when/where to send
 - Meaning of bits
 - => Application programs are the entities that communicate with each other, not the computers or users.
- Important point: For 2 application programs to communicate with each other, one application initiates communication and the other accepts.

Client/Server Model

How 2 application programs make contact



- **Client-server paradigm**: form of communication used by all network applications

Characteristics of a Client

- Arbitrary application program
- Becomes client temporarily
- Can also perform other computations
- Invoked directly by user
- Runs locally on user's computer
- Actively initiates contact with a server
- Contacts one server at a time

Characteristics of a Server

- Special-purpose, privileged program
- Dedicated to providing one service
- Can handle multiple remote clients simultaneously
- Invoked automatically when system boots
- Executes forever
- Needs powerful computer and operating system
- Waits passively for client contact
- Accepts requests from arbitrary clients

Terminology

- Server
 - An executing program that accepts contact over the network
- server-class computer
 - Hardware sufficient to execute a server
- Informally
 - Term “server” often applied to computer

Direction of Data Flow

- Data can flow
 - from client to server only
 - from server to client only
 - in both directions
- Application protocol determines flow
- Typical scenario
 - Client sends request(s)
 - Server sends responses(s)

Server CPU use

- Facts
 - Server operates like other applications
 - uses CPU to execute instructions
 - Performs I/O operations
 - Waiting for data to arrive over a network does not require CPU time
- Consequence
 - Server program uses only CPU when servicing a request

The Socket Interface

- The *Berkeley Sockets API*
 - Originally developed as part of BSD Unix (under gov't grant)
 - BSD = Berkeley Software Distribution
 - API=Application Program Interface
 - Now the most popular API for C/C++ programmers writing applications over TCP/IP
 - Also emulated in other languages: Perl, Tcl/Tk, etc.
 - Also emulated on other operating systems: Windows, etc.

The Socket Interface

- The basic ideas:
 - a *socket* is like a file:
 - you can read/write to/from the network just like you would a file
 - For connection-oriented communication (e.g. TCP)
 - servers (passive open) do **listen** and **accept** operations
 - clients (active open) do **connect** operations
 - both sides can then do **read** and/or **write** (or **send** and **recv**)
 - then each side must **close**
 - There are more details, but those are the most important ideas
 - Connectionless (e.g. UDP): uses **sendto** and **recvfrom**

Sockets And Socket Libraries

- On some other systems, socket procedures are *not* part of the OS
 - instead, they are implemented as a library, linked into the application object code (e.g. a DLL under Windows)
 - Typically, this DLL makes calls to similar procedures that *are* part of the native operating system.
 - This is what the Comer text calls a *socket library*
 - A socket library simulates Berkeley sockets on OS's where the underlying OS networking calls are different from Berkeley sockets

Some definitions

• Data types

int8_t	signed 8-bit integer
int16_t	signed 16-bit integer
int32_t	signed 32-bit integer
uint8_t	unsigned 8-bit integer
uint16_t	unsigned 16-bit integer
uint32_t	unsigned 32-bit integer
u_char	Unsigned 8-bit character
u_short	Unsigned 16-bit integer
u_long	Unsigned 32-bit integer

More Definitions


• Internet Address Structure

```

struct in_addr
{
    in_addr_t s_addr;
};

struct in_addr
{
    u_long s_addr;
};
    
```

Socket address structure

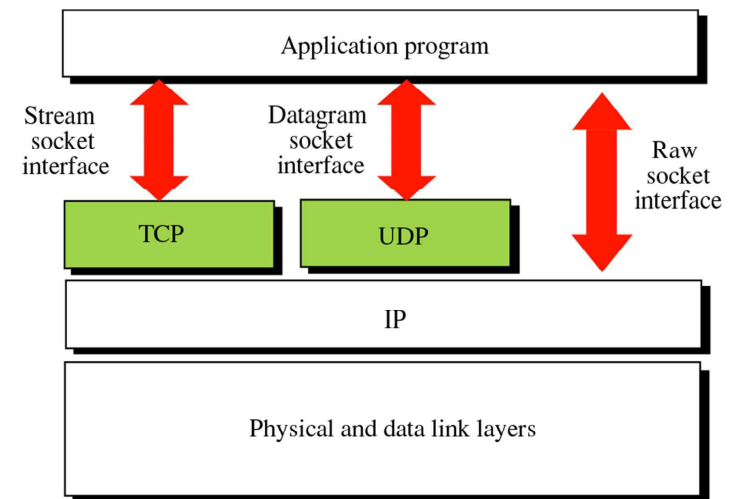


sockaddr_in

```


struct sockaddr_in
{
    u_char sin_len;
    u_short sin_family;
    u_short sin_port;
    struct in_addr sin_addr;
    char sin_zero [8];
};
    
```

Socket Types



Address Transformation

```
int    inet_aton ( const char  *strptr , struct in_addr *addrptr ) ;  
char   *inet_ntoa ( struct in_addr inaddr ) ;
```



Procedures That Implement The Socket API

Creating and Deleting Sockets

- **fd=socket**(*protofamily*, *type*, *protocol*)
Creates a new socket. Returns a file descriptor (*fd*). Must specify:
 - the protocol family (e.g. TCP/IP)
 - the type of service (e.g. STREAM or DGRAM)
 - the protocol (e.g. TCP or UDP)
- **close**(*fd*)
Deletes socket.
For connected STREAM sockets, sends EOF to close connection.

Byte-Manipulation Functions

- In network programming, we often need to initialize a field, copy the contents of one field to another, or compare the contents of two fields.
 - Cannot use string functions (*strcpy*, *strcmp*, ...) which assume null character termination.

```
void *memset ( void *dest , int chr , int len ) ;
```

```
void *memcpy ( void *dest , const void *src , int len ) ;
```

```
int    memcmp ( const void *first , const void *second , int len ) ;
```

Procedures That Implement The Socket API

- **bind**(*fd*)
Used by server to establish port to listen on.
When server has >1 IP addrs, can specify "ANY", or a specific one
- **listen** (*fd*, *queuesize*)
Used by connection-oriented servers only, to put server "on the air"
Queuesize parameter: how many pending connections can be waiting
- **afd = accept** (*lfd*, *caddress*, *caddresslen*)
Used by connection-oriented servers to accept one new connection
 - There must already be a listening socket (*lfd*)
 - Returns *afd*, a new socket for the new connection, and
 - The address of the caller (e.g. for security, log keeping. etc.)

Procedures That Implement The Socket API

How Clients Communicate with Servers

- **connect** (fd, address, saddreslen)

Used by connection-oriented clients to connect to server

- There must already be a socket bound to a connection-oriented service on the fd
- There must already be a listening socket on the server
- You pass in the address (IP address, and port number) of the server.

Used by connectionless clients to specify a “default send to address”

- Subsequent “writes” or “sends” don’t have to specify a destination address
- BUT, there really ISN’T any connection established... this is a bad choice of names!

Procedures That Implement The Socket API

How Clients Communicate with Servers

- **send** (fd, data, length, flags)
- **sendto** (fd, data, length, flags, destaddress, addresslen)
- **sendmsg** (fd, msgstruct, flags)
- **write** (fd, data, length)

Used to send data.

- **send** requires a connection (or for UDP, default send address) be already established
- **sendto** used when we need to specify the dest address (for UDP only)
- **sendmsg** is an alternative version of **sendto** that uses a struct to pass parameters
- **write** is the “normal” write function; can be used with both files and sockets

- **recv** (...) **recvfrom** (...) **recvmsg** (...) **read** (...)

Used to receive data... parameters are similar, but in reverse (destination => source, etc...)

Connectionless Service (UDP)

